

LEANFLOW: A Case Study in Workflow-Driven Lean Autoformalization

Anonymous Authors¹

Abstract

We present and evaluate LEANFLOW, a Lean-specialized LLM-agent system for translating mathematical papers into buildable Lean projects. Recent verifier-in-the-loop systems show that large formal artifacts can be produced, but it remains unclear which runtime mechanisms make document-to-project formalization reliable across model strengths. We study this question on two previously unformalized mathematical papers in number theory and measure theory, using model ablations (Kimi-K2.6 and GPT-5.5) and proof-workflow and toolset ablations; we report task outcome, API calls, input tokens, and output tokens. The full workflow completes both document-level projects with Kimi-K2.6 while no-queue variants hit the 2000-call budget; with GPT-5.5, it lowers model-side token use on both sources. On complementary benchmarks, LEANFLOW reaches 75.7% BEq+ on RLM25, 97% on FATE-H, and solves all five ICML 2026 AI for Math TCS challenge projects in our GPT-5.5 runs.

1. Introduction

Autoformalization is moving from isolated theorem translation toward turning whole mathematical sources into Lean projects. Early language-model formalization work focused on translating individual natural-language statements into proof-assistant syntax (Wu et al., 2022), and many theorem-proving benchmarks still ask the model to prove one supplied formal statement at a time (Zheng et al., 2022; Azerbayev et al., 2023; Jiang et al., 2025). A document-level agent faces a different problem: it must convert mathematical prose, $\text{\TeX}$ structure, references, definitions, and proofs into a buildable Lean project while preserving the intended statements and managing a long sequence of dependent proof obligations. M2F (Wang et al., 2026) and recent

textbook- or project-scale systems (Gloeckle et al., 2026; Hariharan et al., 2026; Tsoukalas et al., 2026) show that verifier-in-the-loop formalization can produce large Lean artifacts; this paper asks which parts of the runtime help agents keep the source claims fixed and finish the resulting proof work.

Motivated by this document-scale setting, LEANFLOW is a Lean-specific runtime that separates mathematical editing from workflow control. It first runs a deterministic source preflight: the runtime resolves the input document, extracts or records source blocks, labels, references, PDFs, figures, and support files, and refuses ambiguous $\text{\TeX}$ roots before the model starts. The formalizer then builds a blueprint, a project-local source map linking source spans to planned Lean declarations, dependencies, proof notes, and recorded scope changes. Before proof search, a statement/source gate checks that the generated declaration types still match the original source claims. During proving, a programmatic workflow manager owns the theorem queue, failed-attempt memory, retry budgets, verification records, logs, and checkpoints. The model proposes edits and helper lemmas, while the manager assigns one target at a time and advances only after external Lean verification accepts the edit.

This separation is useful across the two model settings in our experiments. For Kimi-K2.6 (Moonshot AI, 2026), it turns a long, failure-prone transcript into bounded theorem assignments with explicit context and failure memory. For GPT-5.5, the same structure reduces waste: the model can spend more of its budget on local proof search instead of rediscovering the active theorem, rechecking the whole project, or deciding for itself whether a run is finished.

LEANFLOW also introduces LeanProbe, a cached same-file verifier built on LeanInteract (Poiroux et al., 2025b), for repeated proof-repair checks. LeanProbe gives the proving agent low-latency diagnostics and proof-state feedback while leaving final acceptance to standard Lean/Lake checks. In the sequential same-file benchmark summarized in Table 7, cached checking is roughly 9–14 $\times$ faster than rerunning growing-prefix Lake checks.

Although proof-only benchmarks are a standard way to measure Lean agents, they do not exercise the full document-to-project workflow. In a proof-only setting such as FATE-H (Jiang et al., 2025), the formal statement is already pro-

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

vided and the system only has to replace the proof placeholder. Our main evaluation therefore ablates the components of LEANFLOW on real document-level formalizations: queue management versus free-running proof loops, full Lean tools versus terminal-only interaction, and source-blueprint checkpoints versus direct source-to-code generation. We then report RLM25 statement autoformalization, FATE-H proof-only comparison, and the ICML 2026 AI for Math TCS challenge (AI for Math Workshop, 2026) as complementary measurements.

The two document-level case studies are Frisch and Vaserstein’s *Parametrization of Pythagorean Triples by a Single Triple of Polynomials* (Frisch & Vaserstein, 2007) and Lyons and Zumbrun’s *A Calculus Proof of the Cramer–Wold Theorem* (Lyons & Zumbrun, 2017). The Pythagorean source combines integer-valued polynomial definitions, an obstruction over $\mathbb{Z}[x_1, \dots, x_n]$, explicit rational-polynomial witnesses, and custom plain- $\text{\LaTeX}$ theorem macros. The Cramer–Wold source combines probability measures on Euclidean space, closed half-space values, Crofton reconstruction of average-distance functions, odd-dimensional Laplacian inversion, and an even-to-odd embedding step. These sources are short enough for controlled ablations, but they are still full mathematical papers with nontrivial definitions, representation choices, and proof dependencies. Both require source-aware definitions and reusable Lean infrastructure that the prose leaves implicit.

This paper makes two contributions. First, it describes LEANFLOW as an autonomous workflow for document-to-project formalization and proof repair: the runtime records how source claims become Lean declarations, reviews those declarations before proving, manages the proof queue, and uses LeanProbe for fast local feedback while keeping final acceptance in standard Lean/Lake checks. Second, it evaluates those design choices through ablations under Kimi-K2.6 and GPT-5.5 on document-level case studies, with complementary measurements on RLM25, FATE-H, and the AI for Math TCS challenge.

2. Related Work

Formal theorem proving with language models. Lean 4 (de Moura & Ullrich, 2021) and Mathlib (The mathlib Community, 2020) provide the verification substrate for many recent neural theorem-proving systems. Benchmarks such as miniF2F (Zheng et al., 2022), ProofNet (Azerbayev et al., 2023), PutnamBench (Tsoukalas et al., 2024), FATE (Jiang et al., 2025), FormalMATH (Yu et al., 2025), SorryDB (Letson et al., 2026), and VeriSoftBench (Xin et al., 2026) measure different proof-only or repository-scale capabilities. They provide useful calibration, but most start from an existing formal statement or repository context. LEANFLOW instead studies the runtime needed when the

source begins as mathematical prose, and the output must be a buildable project.

Autoformalization across scales. Early LLM autoformalization work studied natural-language-to-formal translation at theorem scale (Wu et al., 2022); later work improved process supervision and evaluation for Lean statements (Lu et al., 2024; Poiroux et al., 2025c;a). M2F (Wang et al., 2026) and Automatic Textbook Formalization (Gloeckle et al., 2026) formalize long mathematical sources into Lean projects, audit generated statements, and repair proofs under a pinned environment. The Gauss project (Hariharan et al., 2026) completes a sorry-free formalization of the sphere-packing problem in dimension 8 from an existing blueprint and Lean development. AlphaProof Nexus (Tsoukalas et al., 2026) uses an agentic evolutionary framework for formal proof search over a large set of formalized conjectures. Our scope is smaller in document volume but different in emphasis: we isolate queue control, tool routing, statement handoff, and model-side budget under the two model settings in our experiments.

Lean tooling for agents. Retrieval and environment access are central to practical Lean proving. LeanDojo (Yang et al., 2023) studies retrieval-augmented theorem proving, and LeanExplore (Asher, 2025) provides search over Lean declarations. LeanInteract (Poiroux et al., 2025b) exposes a Python interface to Lean execution. LEANFLOW builds on this tool ecosystem but introduces LeanProbe as its own cached same-file verifier for agent loops; final acceptance still comes from standard Lean/Lake verification.

3. Problem Setting

Let D be a mathematical source artifact: a $\text{\LaTeX}$ file, a PDF, or a directory containing a $\text{\LaTeX}$ project. Let E be a pinned Lean environment: a Lean toolchain, a fixed mathematical-library revision such as a Mathlib commit, and the Lake project configuration used for checking. The mathematical result should not depend on this particular snapshot, but the evaluation of builds, diagnostics, imports, and available lemmas is always relative to E . The goal is to produce a Lean project P together with declaration-level provenance. For each generated declaration $d \in \text{Decl}(P)$, the workflow should record

$$\pi(d) \subseteq \text{Span}(D),$$

where $\pi(d)$ is the finite set of source spans, labels, sections, pages, equations, or bibliography references used to justify d .

We separate three checks. Statement type-checking asks whether the generated Lean declaration is well typed under E when theorem-like proof bodies are replaced by explicit `sorry` placeholders. Statement faithfulness asks whether




Figure 1. LEANFLOW document-to-project pipeline. The formalizer emits a well-typed statement skeleton and a durable blueprint; a statement/source gate checks faithfulness; the prover queue then attempts proof closure under verifier and hygiene gates. Redrafts, failed proof attempts, and build failures are persisted in workflow state rather than left in chat history.

that well-typed declaration states the source claim, including object classes, quantifiers, domains, hypotheses, and conclusion. Proof completion asks whether a faithful declaration has a complete proof with no unacceptable placeholders or custom axioms. The final project gate is stricter than a theorem-local check: the requested file or full Lake project must build under E , and a hygiene scan must report no unapproved `sorry`, `admit`, `unsafe`, or hidden axiom.

4. Method

4.1. Workflow Overview

Figure 1 and Algorithm 1 summarize the document-to-project contract. LEANFLOW uses two primary workflows: `formalize`, which turns a source document into source-backed Lean declarations and proof plans, and `prove`, which repairs or completes Lean proofs until the requested file or project compiles successfully. The workflows are deliberately separated. Formalization may introduce definitions, names, imports, and theorem statements, but it stops at a proof skeleton with source-backed `sorry` bodies. Proof repair then treats those declarations as fixed targets: it may edit proof bodies and add local helper lemmas, but it may not modify theorem statements, add axioms, or leave the assigned proof half-solved.

4.2. Source-Backed Skeleton Construction

LEANFLOW begins by resolving the source deterministically. For a `TeX` project directory, deterministic preflight selects a main entry point, follows local includes, records bibliography files, figures, PDFs, and support files, and fails on ambiguous roots before the agent starts. The same pass extracts standard theorem-like environments, custom theorem declarations, plain-`TeX` theorem/proof blocks, labels, references, citations, section structure, and nearby proof environments when available. For PDFs, it records extracted text, metadata, and image inventories through local Poppler tools such as `pdftotext`, `pdfinfo`, and

Algorithm 1 LEANFLOW workflow contract

- 1: Resolve source document D and pinned Lean environment E .
- 2: Extract a preflight manifest: source blocks, labels, references, PDFs, figures, and support files.
- 3: Create a blueprint mapping source spans to planned declarations, dependencies, scope changes, and proof notes.
- 4: Draft Lean declarations with placeholders only for theorem-like proofs; verify that the statement layer type-checks.
- 5: Run a statement/source gate; if faithfulness fails, redraft before proof search.
- 6: **while** the requested file or project still contains assigned proof obligations **do**
- 7: Assign one theorem-like declaration from the queue.
- 8: Use search, proof context, and LeanProbe-backed checks to repair the assigned proof.
- 9: Persist failures, diagnostics, route decisions, and proof checkpoints after each attempt.
- 10: Accept the assignment only after verification, no relevant `sorry`, and an acceptable axiom profile.
- 11: **end while**

`pdfimages`; the model reads this bounded source view through `read.pdf`.

The *formalize* workflow next creates a project-local blueprint before the proof skeleton is considered ready. The blueprint is the durable source map for the run: it lists source statements, planned Lean names, dependencies, source locators, statement-fidelity checks, proof notes, and known scope changes. For each source theorem or lemma, it spells out the parts of the formal type that are easy to drift: object classes, quantifier order, parameter domains, codomains, hypotheses, side conditions, and follow-on equivalences. The generated Lean module must type-check with theorem bodies left as `by sorry`; proof filling is deferred until the statement/source gate has approved the skeleton.

Table 1. Statement/source review checks applied before proof repair.

Check	Rejection signal
Statement type	The type of the generated Lean declaration does not fully match the source statement, so a proof may certify a different claim
Quantifiers	Order, dependency, or implicit parameters changed
Variable types	Variables range over the wrong type
Hypotheses	Added assumption, missing side condition, changed nonemptiness
Conclusion	Weakened equality, image statement, equivalence, or existence
Encoding bridge	A local encoding is used without the definitions or equivalence/evaluation lemmas needed to recover the source statement
Proof notes	Source proof idea missing for a nontrivial theorem
Hygiene	Hidden <code>axiom</code> , <code>unsafe</code> , <code>admit</code> , or unreported <code>sorry</code>

4.3. Statement/Source Gate

The statement/source gate checks the generated skeleton before proof search begins. Compilation only shows that a declaration is well typed; it does not show that the declaration is the theorem stated in the paper. LEANFLOW therefore launches a fresh reviewer context over the blueprint, generated Lean declarations, and original source spans. The reviewer can approve, reject, or request redrafting, but it does not fill proof bodies. If any check in Table 1 fails, the formalizer returns to the skeleton before the *prove* workflow starts; the prover should not receive a weakened, strengthened, or incomparable target. We call the version that passes this gate the *reviewed skeleton*.

4.4. Queue-Managed Proof Repair

The proof workflow is organized around a queue manager rather than a free-form chat loop. The queue manager is a programmatic controller that decides which theorem is active, builds the bounded context shown to the model, runs or routes verification, records failed attempts, and allows the model to advance only when the current assignment is clean. For a file-scoped run, the manager stores the current top-of-queue declaration, the remaining queue, theorem-local failed attempts and diagnostics, retry counters, the latest verification record, and accepted or failed outcomes. The model sees only the current target, a bounded queue horizon, relevant source/proof context, and theorem-local failed attempts. This prevents common long-run failures: losing the active theorem after compaction, retrying a known-bad proof shape, changing a future statement, or advancing because a text search no longer finds `sorry` while Lean still reports open goals.




Figure 2. File-scoped proving loop. Failed checks update theorem-local memory; accepted edits advance to file/Lake verification sweeps.

For project-wide proving, the manager scans files with remaining `sorry`, ranks candidates using declaration-dependency information, local hints/examples, and size estimates, records the plan, and assigns files through the same file-scoped path. After each edit, the proving LLM agent refreshes the structured Lean state and asks the manager whether to continue the theorem, accept the edit, restore a baseline `sorry`, or move to the next target. Figure 2 shows the file-scoped loop.

4.5. LeanProbe and Tool Surface

LeanProbe is a declaration-level verifier for repeated proof attempts in one Lean file. It is a standalone CLI, Python library, and MCP server built on LeanInteract (Poiroux et al., 2025b). At assignment time, LeanProbe runs a *prepare* step for the active declaration: it checks imports, the file header, and all prior declarations, then stores the resulting environment immediately before the target. Each *check* step replaces the target declaration with one candidate proof or helper block and checks it against that cached environment. The cache advances only after a complete declaration replacement is accepted, so failed candidates give diagnostics without changing the environment seen by later obligations.

LeanProbe returns feedback that is directly useful to a proof agent: diagnostics, warnings, `sorry` detection, tactic metadata, goal states, and inline Lean comments at the failing location. A failed candidate can be returned as model-readable Lean with compact comments containing the local proof state and the diagnostic beside the failing tactic, as illustrated in Figure 3.

LeanProbe feedback on a failed candidate

```

theorem add_comm_candidate (x y : Nat) :
  x + y = y + x := by
  /- <feedback>
  -- proof state:
  x y : Nat
  |- x + y = y + x
  </feedback> -/

  rfl
  /- <feedback>
  -- type: error, msg: unsolved goals
  </feedback> -/
    
```

Figure 3. Example LeanProbe feedback embedded into a failed Lean candidate. The proof state is inserted before the failing tactic; the final diagnostic records the open goal.

This information is used for fast local repair inside the queue loop, not for final acceptance. The final gates remain module or project Lake builds and verification that the project is free from unapproved `axiom` declarations and remaining `sorry` placeholders. Table 2 summarizes the tool classes exposed to the model; ablations in Section 5 remove this surface to isolate its effect.

Table 2. Default tool surface exposed to the agent. The manager decides which tool class is relevant for the current blocker.

Tool class	Agent use
Source readers	Inspect TeX/PDF inputs, extracted text, image inventory, and source spans
File editing	Read, search, patch, and write project files under manager checks
Lean verification	Run canonical Lean/Lake checks and LeanProbe cached target checks
Search/context	Retrieve local/Mathlib facts, nearby declarations, and theorem context
Proof exploration	Try bounded proof candidates, request reasoning help, and decompose helpers
Hygiene	Scan for remaining <code>sorry</code> , custom axioms, unsafe constructs, and axiom profiles
Workflow memory	Reattach skills, blueprints, session notes, and theorem-local failed attempts
Terminal fallback	Run explicit shell commands when the bounded tools are insufficient

5. Evaluation

The experiments test whether the LEANFLOW workflow components in Figure 1 improve real formalization and proving runs. Each completed run reports task outcome, API calls, input tokens, and output tokens when frozen workflow logs are available.

5.1. Document-Level Case Studies

The two document-level case studies are selected for mathematical content and area diversity. Frisch and Vaserstein’s Pythagorean-polynomial paper (Frisch & Vaserstein, 2007) requires formalizing the classical $T(a, b, c)$ parametrization, parity and divisibility side conditions, the obstruction to a single integer-coefficient polynomial triple, and the explicit

four-variable triple of integer-valued rational polynomials. Lyons and Zumbrun’s Cramer–Wold paper (Lyons & Zumbrun, 2017) requires formalizing closed half-spaces as measurable sets, equality of measures from matching half-space values, the Crofton step from half-space values to average-distance functions, the odd-dimensional Laplacian recovery argument, and the even-to-odd embedding $\mathbb{R}^{2m} \hookrightarrow \mathbb{R}^{2m+1}$. Section A gives the detailed source-to-Lean decomposition for both case studies. To the best of our knowledge, neither source theorem had a prior proof-assistant formalization.

For the workflow ablations reported in Table 3, each source first runs through the pre-proof stages in Figure 1 and Algorithm 1: source resolution, blueprint construction, statement generation with proof placeholders, and statement/source review. The result is a reviewed skeleton, i.e., a Lean statement layer that type-checks with theorem bodies left as placeholders and has passed the source-faithfulness checks in Table 1. The Pythagorean reviewed skeleton is produced with Kimi-K2.6; the Cramer–Wold reviewed skeleton is produced with GPT-5.5 because the initial representation setup for probability measures, half-spaces, and the analytic inversion target requires more representation choices than the Pythagorean source. The proof runs in Table 3 start from these reviewed skeletons, and there are no human proof edits during the ablated prover runs.

5.2. Workflow Ablations on Real Sources

The evaluation is based on ablation: it compares the full LEANFLOW against variants that remove the queue manager and/or the specialized Lean tools. For each model, we run the full 2×2 design: queue versus no queue, and full tool surface versus terminal-only interaction. The corresponding Table 3 is organized by source so that the final analysis can compare completion behavior and model-side budget within the same mathematical project. All document-level proving runs have a maximum budget of 2000 prover-agent API calls, where one call is one model turn in the proof-repair loop. In the no-queue condition, an outer runner still keeps the prover running until the project succeeds or the call budget expires, but it disables the assignment and memory behavior in Algorithm 1, Step 6: the model is not restricted to a single top-of-queue theorem, and theorem-local failed-attempt history is not used to gate the next target. Source preflight, the statement/source gate, final Lean/Lake verification, and hygiene checks are held fixed across conditions.

Table 3 shows that queue management changes the outcome for Kimi-K2.6 on these document-level sources: both no-queue variants exhaust the 2000-call budget on both case studies, while full LEANFLOW succeeds in 1043 calls on Pythagorean triples and 1278 calls on Cramer–Wold. The tool surface matters most on the analytic case study, where Kimi-K2.6 with a queue but terminal-only interaction still

Table 3. Long-form workflow ablation on the two document-level case studies. Workflow abbreviations encode queue and tool configuration: Full = queue with full tools; NoQ+Tools = no queue with full tools; Q+CLI = queue with terminal-only interaction; NoQ+CLI = no queue with terminal-only interaction. Calls are prover-agent API calls with a 2000-call cap.

Source	Model	Workflow	Outcome	Calls	In tok.	Out tok.
Pythagorean	Kimi-K2.6	Full	success	1043	46.9M	514.1K
Pythagorean	Kimi-K2.6	NoQ+Tools	failure	2000	160.2M	682.0K
Pythagorean	Kimi-K2.6	Q+CLI	success	1561	91.3M	537.8K
Pythagorean	Kimi-K2.6	NoQ+CLI	failure	2000	157.5M	615.5K
Cramer–Wold	Kimi-K2.6	Full	success	1278	66.0M	684.1K
Cramer–Wold	Kimi-K2.6	NoQ+Tools	failure	2000	127.8M	583.2K
Cramer–Wold	Kimi-K2.6	Q+CLI	failure	2000	113.1M	579.0K
Cramer–Wold	Kimi-K2.6	NoQ+CLI	failure	2000	167.0M	496.4K
Pythagorean	GPT-5.5	Full	success	258	14.3M	190.5K
Pythagorean	GPT-5.5	NoQ+Tools	success	568	45.0M	240.3K
Pythagorean	GPT-5.5	Q+CLI	success	277	14.3M	177.0K
Pythagorean	GPT-5.5	NoQ+CLI	success	984	89.9M	340.7K
Cramer–Wold	GPT-5.5	Full	success	495	25.9M	305.5K
Cramer–Wold	GPT-5.5	NoQ+Tools	success	448	32.8M	259.7K
Cramer–Wold	GPT-5.5	Q+CLI	success	959	59.5M	518.0K
Cramer–Wold	GPT-5.5	NoQ+CLI	success	476	49.7M	281.6K

Table 4. RLM25 statement-autoformalization results. Type-check reports the share of generated statements accepted by Lean; BEq and BEq+ follow the benchmark protocol.

Model	Workflow	Type-check	BEq	BEq+	Calls	In tok.	Out tok.
GPT-5.5	LEANFLOW	81.2%	70.8%	75.7%	3541	203.2M	2.1M
GPT-5.5	terminal-only agent	80.6%	68.1%	72.9%	4053	236.6M	2.3M

fails on Cramer–Wold. For GPT-5.5, all document-level variants succeed, so the effect is mostly budget use rather than binary success; full LEANFLOW has the lowest or tied-lowest input-token cost on both sources, although the Cramer–Wold no-queue/full-tools run uses slightly fewer calls at higher input-token cost.

5.3. RLM25 Autoformalization

We use RLM25 (Poiroux et al., 2025a;c) to test theorem-statement autoformalization for research mathematics under the same model/tool distinction. RLM25 is a collection of informal theorem statements drawn from research-level formalization projects; the system must synthesize the formal Lean statement, and the reference formal statement is used only for evaluation through BEq and BEq+. Because RLM25 contains independent theorem statements rather than complete source documents, we report it separately from the document-to-project case studies. Table 4 summarizes the completed GPT-5.5 runs. For GPT-5.5, LEANFLOW improves BEq+ from 72.9% to 75.7% while reducing calls and token use. The gap between type-checking and semantic-equivalence metrics reinforces why statement/source checks cannot be replaced by compilation alone: a well-typed declaration still needs an external faithfulness check against the intended mathematical claim.

Table 5. FATE-H proof-only comparison. The M2F row reports the published proof-success result; calls and tokens are available for our run only.

Model	Solved	Calls	In tok.	Out tok.
LEANFLOW (GPT-5.5)	97%	11729	932.3M	5.4M
M2F (GPT5.2 Pro)	96%	-	-	-

5.4. Proof-Only and Challenge Tasks

Proof-only tasks isolate theorem proving once formal statements are already provided. We first use FATE-H as a comparison point against M2F’s proof-only result, as shown in Table 5.

LEANFLOW solves 97% of FATE-H, close to the 96% proof-only result reported for M2F, but the run still requires 11,729 prover-agent calls and 932.3M input tokens. This supports using FATE-H as a hard proof-only comparison while keeping the main claims centered on document-level workflow behavior.

The ICML 2026 AI for Math Track 2 TCS proving challenge tests project-level algorithmic proof tasks in graph theory, combinatorics, computational complexity, and algorithm correctness using CSLib and a comparator-backed pass-rate metric (AI for Math Workshop, 2026). The five challenge projects cover binary heap operations, a heap-backed Dijkstra proof, Kruskal minimum spanning trees, segment-tree construction and range queries, and treap invariant/runtime obligations. We run all five projects with GPT-5.5 under three conditions: full LEANFLOW, no queue with LeanProbe still available, and no queue with only terminal Lean/Lake interaction. For Kimi-K2.6, we additionally run the Kruskal project with full LEANFLOW and with nei-

ther queue management nor the specialized tool surface.

Table 6 shows that all GPT-5.5 variants solve all five TCS projects, so these rows do not isolate a single decisive workflow component. Still, the results are consistent with LeanProbe being useful as a feedback surface: workflow logs show that GPT-5.5 used LeanProbe in 37% of recorded agent steps in TCS runs where the tool was available, and NoQ+Probe uses fewer calls and lower input tokens than NoQ+CLI on four of five projects. This aligns with the tool design in Section 4.5, where failed candidates can return local compiler diagnostics and proof-state feedback before the next proof attempt. For Kimi-K2.6, the tested Kruskal run succeeds with the full workflow, while the no-queue terminal-only variant reaches the 2000-call budget.

5.5. LeanProbe Utility

LeanProbe is evaluated separately because it changes verifier latency rather than model capability. We compare cached LeanProbe checks against terminal `lake env lean` checks on two workloads that match the proof loop in Figure 2. The repeated-target workload measures many candidate replacements for one active declaration after a single prepare step. The sequential workload measures a queue-style file run: partial candidates are rejected against the current cached environment, while accepted declarations advance the environment for later targets.

Table 7 explains why LEANFLOW checks candidate edits frequently instead of waiting for large repair batches. The prepare step is paid once before a target, and then each can-

didate proof replacement is checked in milliseconds to tens of milliseconds. Sequential checks also match the queue design: the environment advances only after a complete declaration succeeds, so failed candidates are diagnosed without changing the context for later obligations. The detailed LeanProbe experimental setup and per-file results are reported in Section B.

5.6. Statement-Skeleton Case Study

We use a qualitative Pythagorean-polynomial comparison to check whether the source-to-statement stages in Algorithm 1 prevent statement drift before proof search. The no-review condition skips the blueprint and statement/source review stages in Steps 3 and 5: after seeing the source, the model is asked to produce Lean translations directly, without a recorded source-to-declaration map or an independent faithfulness check before proving. The no-review draft and the reviewed skeleton both build, but they formalize different mathematical objects. The no-review draft would support a less source-faithful function-level formalization of the result: it defines the main witnesses as rational-valued functions on integer inputs. The reviewed skeleton is the source-faithful formalization: it defines the witnesses as rational multivariate polynomials and separately proves integer-valuedness, matching the paper’s theorem. The detailed declaration-level contrast is reported in the Appendix A (Table 8).

Table 6. AI4Math TCS challenge outcomes by project and workflow condition. Workflow abbreviations encode queue and tool configuration: Full = queue with full tools; NoQ+Probe = no queue with LeanProbe available; NoQ+CLI = no queue with terminal-only Lean/Lake interaction. Calls are prover-agent API calls.

Project	Model	Workflow	Outcome	Calls	In tok.	Out tok.
Binary heap	GPT-5.5	Full	success	54	2.86M	49.8K
Binary heap	GPT-5.5	NoQ+Probe	success	27	741.3K	7.1K
Binary heap	GPT-5.5	NoQ+CLI	success	72	5.03M	71.1K
Dijkstra	GPT-5.5	Full	success	61	4.04M	75.1K
Dijkstra	GPT-5.5	NoQ+Probe	success	20	770.2K	26.0K
Dijkstra	GPT-5.5	NoQ+CLI	success	29	1.15M	23.9K
Segment tree	GPT-5.5	Full	success	62	4.01M	35.0K
Segment tree	GPT-5.5	NoQ+Probe	success	85	7.01M	42.2K
Segment tree	GPT-5.5	NoQ+CLI	success	73	4.67M	45.4K
Treap	GPT-5.5	Full	success	142	10.88M	85.1K
Treap	GPT-5.5	NoQ+Probe	success	144	10.80M	95.6K
Treap	GPT-5.5	NoQ+CLI	success	188	14.10M	116.8K
Kruskal	GPT-5.5	Full	success	47	1.75M	16.2K
Kruskal	GPT-5.5	NoQ+Probe	success	37	1.56M	17.7K
Kruskal	GPT-5.5	NoQ+CLI	success	53	2.94M	26.8K
Kruskal	Kimi-K2.6	Full	success	942	107.9M	3.1M
Kruskal	Kimi-K2.6	NoQ+CLI	failure	2000	189.1M	930.4K

Table 7. LeanProbe latency benchmarks. Repeated-target times are average ranges in seconds after one prepare step; sequential rows report total seconds per compact file with five declarations and ten partial/full candidate scenarios.

Workload	OS	Full-file Lake	Cached check	Workload	OS	Growing-prefix Lake	Cached total	Speedup
TCS	macOS	2.082–2.617s	0.031–0.049s	Compact files	macOS	36.474–67.992s	3.789–4.775s	9.63–14.24×
TCS	Linux	1.495–1.886s	0.032–0.054s	Compact files	Linux	22.593–23.186s	2.301–2.547s	9.05–9.82×

6. Threats to Validity

The evaluation has several limitations. First, project buildability is necessary but not sufficient: a project can build with a faithful theorem, a weakened theorem, or a theorem whose representation silently changed. We mitigate this through the statement/source gate, but the document-level skeletons still rely on expert judgment before proof ablations begin; source spans and blueprint decisions should therefore be released with the artifacts so disagreements are inspectable. Second, two document-level case studies cannot establish broad coverage across mathematics, and the Cramer–Wold skeleton was initialized with GPT-5.5 rather than Kimi-K2.6 because of its representation choices. Third, model comparisons are sensitive to model version, serving stack, context limit, retry policy, and tool latency. Calls and token counts are useful proxies for model-side effort, but they are not a full measure of wall-clock cost or human engineering cost.

7. Conclusion

LEANFLOW treats autoformalization as an evolving Lean project with durable workflow state, not as independent theorem prompts. The central design choice is to keep project control outside the model: the runtime records source-to-Lean decisions, reviews statements before proof search, assigns proofs one at a time, checks candidate edits quickly, and accepts only builds that pass the verifier and contain no leftover placeholders or unapproved axioms. The current results show that this structure lets Kimi-K2.6 finish document-level sources that the no-queue variants do not finish within budget, helps GPT-5.5 spend less model budget, and separates statement synthesis, proof-only proving, and algorithmic proof workloads.

Impact Statement

This paper presents work whose goal is to make Lean autoformalization more reliable, inspectable, and reproducible. Potential positive impacts include reducing the engineering burden of proof repair, improving the auditability of machine-assisted formalization, and making document-level formal verification workflows easier to study. Potential risks include over-reliance on generated formal artifacts, mistaking buildability for source faithfulness, and unequal access to the compute and model resources needed for long formal-

ization runs. LEANFLOW is designed to mitigate these risks by recording source-to-Lean decisions, separating statement review from proof search, preserving workflow logs and checkpoints, and accepting final artifacts only after external Lean verification with no leftover placeholders or unapproved axioms. The system is intended to assist expert formalizers rather than replace human mathematical judgment.

References

- AI for Math Workshop. ICML 2026 AI for math workshop challenge tracks. <https://ai4math2026.github.io/>, 2026. Workshop challenge description, including Track 2 theoretical computer science proving in Lean.
- Asher, J. LeanExplore: A search engine for Lean 4 declarations, 2025. URL <https://arxiv.org/abs/2506.11085>.
- Azerbaiyev, Z., Piotrowski, B., Schoelkopf, H. W., Ayers, E. W., Radev, D., and Avigad, J. ProofNet: Autoformalizing and formally proving undergraduate-level mathematics, 2023. URL <https://arxiv.org/abs/2302.12433>.
- de Moura, L. and Ullrich, S. The Lean 4 theorem prover and programming language. In *Automated Deduction – CADE 28*, pp. 625–635. Springer, 2021. doi: 10.1007/978-3-030-79876-5_37.
- Frisch, S. and Vaserstein, L. Parametrization of Pythagorean triples by a single triple of polynomials, 2007. URL <https://arxiv.org/abs/0706.0290>.
- Gloeckle, F., Rammal, A., Arnal, C., Munos, R., Cabannes, V., Synnaeve, G., and Hayat, A. Automatic textbook formalization, 2026. URL <https://arxiv.org/abs/2604.03071>.
- Hariharan, S., Birkbeck, C., Lee, S., Ma, H. K. G., Mehta, B., Poiroux, A., and Viazovska, M. A milestone in formalization: The sphere packing problem in dimension 8, 2026. URL <https://arxiv.org/abs/2604.23468>.
- Jiang, J., He, W., Wang, Y., Gao, G., Hu, Y., Wang, J., Guan, N., Wu, P., Dai, C., Xiao, L., and Dong, B. FATE: A

- formal benchmark series for frontier algebra of multiple difficulty levels, 2025. URL <https://arxiv.org/abs/2511.02872>.
- Letson, A., Sarra, L., Poiroux, A., Dressler, O., Lezeau, P., Aranha, D., Pu, F., Hill, A., Hidalgo, M. C., Berman, J., Tsoukalas, G., and Taelman, L. SorryDB: Can AI provers complete real-world Lean theorems?, 2026. URL <https://arxiv.org/abs/2603.02668>.
- Lu, J., Wan, Y., Liu, Z., Huang, Y., Xiong, J., Liu, C., Shen, J., Jin, H., Zhang, J., Wang, H., Yang, Z., Tang, J., and Guo, Z. Process-driven autoformalization in Lean 4, 2024. URL <https://arxiv.org/abs/2406.01940>.
- Lyons, R. and Zumbur, K. A calculus proof of the Cramér–Wold theorem, 2017. URL <https://arxiv.org/abs/1607.03206>.
- Moonshot AI. Kimi-K2.6: Open-weights native multimodal agentic model. <https://huggingface.co/moonshotai/Kimi-K2.6>, 2026. Hugging Face model card for moonshotai/Kimi-K2.6.
- Poiroux, A., Bosselut, A., and Kuncak, V. RLMEval: Evaluating research-level neural theorem proving, 2025a. URL <https://arxiv.org/abs/2510.25427>.
- Poiroux, A., Kuncak, V., and Bosselut, A. LeanInteract: A python interface for Lean 4. <https://github.com/augustepoiroux/LeanInteract>, 2025b.
- Poiroux, A., Weiss, G., Kunčak, V., and Bosselut, A. Reliable evaluation and benchmarks for statement autoformalization. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 17947–17969, Suzhou, China, 2025c. Association for Computational Linguistics. doi: 10.18653/v1/2025.emnlp-main.907. URL <https://aclanthology.org/2025.emnlp-main.907/>.
- The mathlib Community. The Lean mathematical library. In *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs*, pp. 367–381. Association for Computing Machinery, 2020. doi: 10.1145/3372885.3373824.
- The Physlib community. Physlib: The Lean physics library. <https://github.com/leanprover-community/physlib>, 2024. Formerly PhysLean/HepLean.
- Tsoukalas, G., Lee, J., Jennings, J., Xin, J., Ding, M., Jennings, M., Thakur, A., and Chaudhuri, S. Putnam-Bench: Evaluating neural theorem-provers on the putnam mathematical competition, 2024. URL <https://arxiv.org/abs/2407.11214>.
- Tsoukalas, G., Kovsharov, A., Shirobokov, S., Surina, A., Firsching, M., Bérczi, G., Ruiz, F. J. R., Suggala, A., Wagner, A. Z., Wieser, E., Yu, L., Huang, A., Horváth, M. Z., Ferraiuolo, A., Michalewski, H., Grosu, C., Hubert, T., Balog, M., Kohli, P., and Chaudhuri, S. Advancing mathematics research with ai-driven formal proof search, 2026. URL <https://arxiv.org/abs/2605.22763>.
- Wang, Z., Ma, W., Ming, Z., Zhang, G., Yuan, K., and Wen, Z. M2F: Automated formalization of mathematical literature at scale, 2026. URL <https://arxiv.org/abs/2602.17016>.
- Wu, Y., Jiang, A. Q., Li, W., Rabe, M. N., Staats, C., Jamnik, M., and Szegedy, C. Autoformalization with large language models, 2022. URL <https://arxiv.org/abs/2205.12615>.
- Xin, Y., Chen, Q., Durrett, G., and Dillig, I. VeriSoft-Bench: Repository-scale formal verification benchmarks for Lean, 2026. URL <https://arxiv.org/abs/2602.18307>.
- Yang, K., Swope, A. M., Gu, A., Chalamala, R., Song, P., Yu, S., Godil, S., Prenger, R., and Anandkumar, A. LeanDojo: Theorem proving with retrieval-augmented language models, 2023. URL <https://arxiv.org/abs/2306.15626>.
- Yu, Z., Peng, R., Ding, K., Li, Y., Peng, Z., Liu, M., Zhang, Y., Zheng, Y., Xin, H., Huang, W., Wen, Y., and Liu, W. FormalMATH: Benchmarking formal mathematical reasoning of large language models, 2025. URL <https://arxiv.org/abs/2505.02735>.
- Zheng, K., Han, J. M., and Polu, S. MiniF2F: A cross-system benchmark for formal olympiad-level mathematics. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=9ZPegFuFTFv>.

A. Document Case Study Details

Pythagorean-polynomial source. The Frisch–Vaserstein source first defines Pythagorean triples, polynomial parametrization over $\mathbb{Z}[x_1, \dots, x_n]$, and integer-valued parametrization over $\text{Int}(\mathbb{Z}^n) \subseteq \mathbb{Q}[x_1, \dots, x_n]$. The formalization therefore introduces both integer-coefficient and rational-polynomial representations before stating the main claims. The source-level proof obligations are: the standard two-family parametrization of triples, the obstruction to a single integer-coefficient triple, the explicit four-variable integer-valued witness, the positive-triple variant, and the four-square unrestricted-parameter variant. The obstruction proof is algebraic: it uses the UFD structure of $\mathbb{Z}[x_1, \dots, x_n]$, a gcd decomposition, parity facts, and the examples $(3, 4, 5)$ and $(4, 3, 5)$ to force a contradiction. The construction proof uses the source map

$$T(a, b, c) = (c(a^2 - b^2)/2, cab, c(a^2 + b^2)/2)$$

and the parity condition that makes all three coordinates integral. The reviewed Lean layout separates these concerns into basic definitions, source handoff lemmas, obstruction lemmas, integer-valued construction lemmas, positive-triple lemmas, and explanatory material for cited background. This layout also records source claims that are easy to lose in a proof-only draft: the finite-cover theorem for integer-valued parametrizations, the displayed binomial-polynomial identity, and the non-UFD motivation for $\text{Int}(\mathbb{Z})$.

Table 8. Concrete declaration-level gap in the Pythagorean statement skeleton. The no-review draft is generated directly from the source; the reviewed skeleton is the version accepted by the statement/source review conditions described in Table 1.

Source requirement	No-review draft	Reviewed skeleton
Integer-valued polynomial objects	<code>IntegerValued4</code> is a predicate definition on functions $\mathbb{Z} \rightarrow \mathbb{Z} \rightarrow \mathbb{Z} \rightarrow \mathbb{Z} \rightarrow \mathbb{Q}$; integer-valuedness is checked only after evaluation at integer inputs.	<code>RatPoly n</code> is a type abbreviation for rational multivariate polynomials; <code>IsIntValued</code> is a predicate definition, and <code>IntValuedSubring n</code> is a subring definition.
Displayed witnesses are polynomials	<code>f_param</code> , <code>g_param</code> , and <code>h_param</code> are function definitions of four integer arguments, so the displayed source witnesses are encoded only by their evaluations.	<code>f_param</code> , <code>g_param</code> , and <code>h_param</code> are polynomial definitions of type <code>RatPoly 4</code> ; evaluation at integer tuples is handled separately by <code>ratPolyEval</code> .
Main theorem type	<code>integer_valued_</code> is a theorem proving integer-parametrization valuedness and image equality for the function model.	<code>exists_int_valued_</code> is a theorem with conclusion <code>IntValuedParametrizes ...</code> , so it states existence of integer-valued rational polynomials that parametrize the triples.
Background claims retained	The direct translation omits theorem or definition declarations for the finite-cover result and the integer-valued-polynomial ring background.	The reviewed skeleton keeps theorem declarations for the finite-cover result and non-UFD claim, plus polynomial definitions and theorems for the binomial-polynomial identity.

Cramer–Wold source. The Lyons–Zumbrun source states that Borel probability measures on $\mathbb{R}^n$ are determined by their values on closed half-spaces. The formalization represents $\mathbb{R}^n$ as Lean Euclidean space, defines closed half-spaces by a normal vector and threshold, and records half-space value functions and the average-distance function

$$f_\mu(y) = \int (\|y - x\| - \|x\|) d\mu(x).$$

The first analytic proof block packages the Crofton/Fubini argument showing that half-space values determine f_μ . The second proof block packages the odd-dimensional inversion step: after reindexing dimensions as $2m + 1$, the needed kernel identity is that an $(m + 1)$ -fold Laplacian of the norm kernel recovers a nonzero multiple of the point mass at the origin. In Lean-facing terms, this required a tempered-distribution norm kernel, Fourier/Laplacian algebra for iterated Laplacians, transport between the Euclidean model used by PhysLean (The Physlib community, 2024) and the project’s odd-dimensional space, and a bridge from PhysLean’s real distribution API to Mathlib’s complex tempered distributions. The even-dimensional Cramer–Wold theorem is then reduced to the odd-dimensional one by embedding $\mathbb{R}^{2m}$ into $\mathbb{R}^{2m+1}$ and transporting half-space values along that embedding. The reviewed skeleton keeps the Crofton and inversion blocks as named proof targets rather than hiding them inside the final theorem; this makes it clear which missing analytic facts are mathematical infrastructure, not model-generated assumptions.

B. LeanProbe Benchmark Results

This appendix records the detailed LeanProbe benchmark rows behind Table 7. The snapshot was refreshed on May 13, 2026 with Lean 4.30.0-rc2. These rows are verifier-latency measurements, not model-performance measurements: they isolate the cost of checking candidate Lean edits under different reuse policies. The benchmark suite has two parts: grouped repeated-target checks, which stress many alternatives for one active declaration, and sequential same-file checks, which stress queue-style advancement through a file. All runs use one measured run per target/file, no benchmark warmups, warm Lake caches from prior validation, and a fresh-server baseline. LeanProbe remains a feedback surface; final file or project acceptance still uses standard Lean/Lake verification.

Grouped repeated-target benchmark. The repeated-target workload measures the agent pattern of trying several complete replacements for one active declaration. For each target, the benchmark writes a temporary full file and times terminal `lake env lean`; then it starts LeanProbe, prepares the environment immediately before the target, checks the replacement against that cached environment, optionally requests proof-state feedback, and finally repeats the check with a fresh LeanProbe/LeanInteract server to measure the no-cache baseline. The compact groups are Mathlib-facing examples: `analysis_real` covers elementary real-analysis facts, `algebra_order` covers ordered algebra and inequalities, `sets_functions` covers set/function image and preimage reasoning, and `number_theory_nat` covers elementary natural-number arithmetic. The TCS groups are longer challenge-style extracts: `tcs_binary_heap` contains heap operations such as `heapify`, `extract_min`, `insert`, `merge`, and `remove`; `tcs_treap_analysis` contains the two probability-sum obligations used in the treap analysis; and `tcs_weighted_graph` contains weighted-graph helper declarations through the `Sym2order` prefix. In Table 9, *Targets* is the number of active declarations in the group, *Lake full* is full-file terminal verification, *Prepare* is the one-time cached-environment setup, *Cached* is a target replacement check after prepare, *Feedback* includes diagnostic/proof-state metadata, *Fresh* is the same check with a fresh server, and *Fresh/cached* is the fresh-check time divided by the cached-check time.

Sequential same-file benchmark. The sequential workload models a queue-managed file run rather than repeated attempts on one target. Each compact file has five declarations and ten scenarios: for every targetable declaration, the benchmark checks a partial declaration containing `sorry` and then the complete declaration. LeanProbe reports the partial scenario and detects the `sorry`, but it advances the cached environment only after the complete declaration succeeds. The Lake baselines rerun terminal checks for every scenario, either on a growing prefix containing accepted prior declarations or on a full temporary file with only the current declaration replaced. In Table 10, *Lake prefix* and *Lake full* are these two terminal baselines, *Probe cached* is one warm LeanProbe/LeanInteract server walking the whole file, *Probe fresh* restarts LeanProbe for every scenario, and the three speedup columns divide the corresponding baseline total by the cached LeanProbe total.

Table 9. LeanProbe repeated-target grouped results. Times are averages in seconds; fresh/cached is the ratio between a fresh LeanProbe server and a cached check.

Platform	Group	Targets	Lake full	Prepare	Cached	Feedback	Fresh	Fresh/cached
macOS	<code>analysis_real</code>	5	3.893	6.024	0.039	0.022	4.014	139.7×
macOS	<code>algebra_order</code>	5	3.900	3.683	0.048	0.039	3.987	106.7×
macOS	<code>sets_functions</code>	5	3.708	3.502	0.008	0.007	3.766	454.7×
macOS	<code>number_theory_nat</code>	5	3.731	3.478	0.011	0.006	3.776	420.0×
Linux	<code>analysis_real</code>	5	2.276	2.315	0.025	0.024	2.412	103.2×
Linux	<code>algebra_order</code>	5	2.301	2.317	0.046	0.043	2.516	78.2×
Linux	<code>sets_functions</code>	5	2.233	2.257	0.011	0.009	2.383	245.8×
Linux	<code>number_theory_nat</code>	5	2.199	2.217	0.009	0.008	2.390	322.0×
macOS	<code>tcs_binary_heap</code>	9	2.576	2.931	0.049	0.042	2.589	155.9×
macOS	<code>tcs_treap_analysis</code>	2	2.082	2.219	0.034	0.034	2.181	77.8×
macOS	<code>tcs_weighted_graph</code>	9	2.617	2.461	0.031	0.028	2.603	194.9×
Linux	<code>tcs_binary_heap</code>	9	1.886	1.807	0.054	0.051	1.877	103.2×
Linux	<code>tcs_treap_analysis</code>	2	1.495	1.441	0.036	0.040	1.560	53.1×
Linux	<code>tcs_weighted_graph</code>	9	1.771	1.683	0.032	0.034	1.761	127.5×

Table 10. LeanProbe sequential same-file results. Each file has five declarations and ten partial/full scenarios. Times are total seconds for the file.

Platform	File	Decls.	Scen.	Lake prefix	Lake full	Probe cached	Probe fresh	vs prefix	vs full	vs fresh
macOS	analysis_real	5	10	67.992	40.216	4.775	44.699	14.24×	8.42×	9.36×
macOS	algebra_order	5	10	46.870	48.163	4.339	40.953	10.80×	11.10×	9.44×
macOS	sets_functions	5	10	45.421	42.237	3.916	37.797	11.60×	10.79×	9.65×
macOS	number_theory_nat	5	10	36.474	36.648	3.789	36.736	9.63×	9.67×	9.70×
Linux	analysis_real	5	10	22.765	22.958	2.515	24.890	9.05×	9.13×	9.90×
Linux	algebra_order	5	10	23.186	23.489	2.547	25.147	9.10×	9.22×	9.87×
Linux	sets_functions	5	10	22.679	22.567	2.384	24.195	9.51×	9.47×	10.15×
Linux	number_theory_nat	5	10	22.593	22.829	2.301	24.086	9.82×	9.92×	10.47×

Repeated-target rows isolate the cost of checking many replacements for one active declaration after a single prepare step. For n complete replacement attempts on the same target, the comparison is

$$\text{Probe total}(n) = \text{Prepare} + n \cdot \text{Cached}, \quad \text{Lake total}(n) = n \cdot \text{Lake full}.$$

The prepare cost is therefore amortized as soon as a target receives several candidate repairs: after setup, each additional candidate costs the cached-check time instead of another full-file `lake env lean run`. The grouped rows report this steady-state behavior directly through *Cached*, *Fresh*, and *Fresh/cached*; the sequential rows report the end-to-end cost when successful declarations advance the environment. Sequential rows model queue-style proving in one file: partial `sorry` scenarios are diagnosed but not cached, while complete accepted declarations advance the environment for later targets. This matches the proof-loop policy used in our runs: keep the LeanProbe server warm, prepare once before a target when several repairs are likely, and advance the cached environment only after a complete declaration has been accepted.